

## 6.2.a ConvNet: CIFGAR10 image classifier

```
In [1]: import os
        from google.colab import drive
        drive.mount('/content/drive', force_remount = True)
        os.chdir('/content/drive/My Drive/DSC650/assignment06')
        !pwd
```

Mounted at /content/drive  
/content/drive/My Drive/DSC650/assignment06

```
In [2]: import numpy as np
        import pandas as pd
        import matplotlib.pyplot as plt
        import pickle

        from keras import layers, models
        from keras.datasets import cifar10
        from keras.models import Sequential
        from keras.layers import Dense, Dropout, Activation
        from keras.utils import np_utils, to_categorical
        from keras.optimizers import SGD

        from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau, ModelCheckpoint
```

```
In [3]: (trainX, trainy), (testX, testy) = cifar10.load_data()
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>  
170498071/170498071 [=====] - 5s 0us/step

```
In [4]: # get the size of the data sets
        print(f'train_images: {trainX.shape}')
        print(f'test_images: {testX.shape}')
        print(f'train_labels: {trainy.shape}')
        print(f'test_labels: {testy.shape}')
```

train\_images: (50000, 32, 32, 3)  
test\_images: (10000, 32, 32, 3)  
train\_labels: (50000, 1)  
test\_labels: (10000, 1)


```
In [5]: # Assignment classes for visualization
        cifar10_classes = ['airplane', 'automobile', 'bird', 'cat',
                           'deer', 'frog', 'horse', 'ship', 'truck']
```

Visualize sample images

```
In [6]: fig, ax = plt.subplots(5, 5)
        k = 0

        for i in range(5):
            for j in range(5):
                ax[i][j].imshow(trainX[k], aspect = 'auto')
                k += 1

        plt.show()
```



```
In [7]: # normalize datasets
train_images = trainX.astype('float32') / 255.0
test_images = testX.astype('float32') / 255.0
```

```
In [8]: # convert labels to numeric
train_labels = to_categorical(trainy)
test_labels = to_categorical(testy)
```

Split training data into training and validation datasets

```
In [9]: x_val = train_images[:10000]
partial_x_train = train_images[10000:]

y_val = train_labels[:10000]
partial_y_train = train_labels[10000:]
```

```
In [10]: # get the size of the data sets
print(f'x_val: {x_val.shape}')
print(f'y_val: {y_val.shape}')
print(f'partial_x_train: {partial_x_train.shape}')
print(f'partial_y_train: {partial_y_train.shape}')
```

```
x_val: (10000, 32, 32, 3)
y_val: (10000, 10)
partial_x_train: (40000, 32, 32, 3)
partial_y_train: (40000, 10)
```

Build the Model

```
In [11]: # Instantiate a convnet
model = Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu', kernel_initializer='he_uniform'))
model.add(layers.Conv2D(32, (3, 3), activation='relu', kernel_initializer='he_uniform'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform'))
model.add(layers.Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform'))
model.add(layers.Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Flatten())
```

```
model.add(Dense(128, activation='relu', kernel_initializer='he_uniform'))
model.add(Dense(10, activation='softmax'))
```

In [12]: `model.summary()`

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	896
conv2d_1 (Conv2D)	(None, 32, 32, 32)	9248
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_2 (Conv2D)	(None, 16, 16, 64)	18496
conv2d_3 (Conv2D)	(None, 16, 16, 64)	36928
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 64)	0
conv2d_4 (Conv2D)	(None, 8, 8, 128)	73856
conv2d_5 (Conv2D)	(None, 8, 8, 128)	147584
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 128)	0
flatten (Flatten)	(None, 2048)	0
dense (Dense)	(None, 128)	262272
dense_1 (Dense)	(None, 10)	1290
Total params: 550,570		
Trainable params: 550,570		
Non-trainable params: 0		

Compile the Model

In [13]: `opt = SGD(lr = 0.001, momentum = 0.9)`

```
model.compile(loss = 'categorical_crossentropy',
              optimizer = opt,
              metrics = ['accuracy'])
```

/usr/local/lib/python3.8/dist-packages/keras/optimizers/optimizer\_v2/gradient\_descent.py:108: UserWarning: The `lr` argument is deprecated, use `learning\_rate` instead.  
super(SGD, self).\_\_init\_\_(name, \*\*kwargs)

Train the model

In [15]: `history = model.fit(partial_x_train, partial_y_train,`  
`epochs = 30, batch_size = 64,`  
`validation_data=(x_val, y_val), verbose = 1)`

Epoch 1/30  
625/625 [=====] - 5s 8ms/step - loss: 0.0087 - accuracy: 0.9982 - val\_loss: 2.1163 - val\_accuracy: 0.7056  
Epoch 2/30  
625/625 [=====] - 5s 8ms/step - loss: 0.0037 - accuracy: 0.9995 - val\_loss: 2.1002 - val\_accuracy: 0.7178  
Epoch 3/30  
625/625 [=====] - 5s 8ms/step - loss: 8.5766e-04 - accuracy: 1.0000 - val\_loss: 2.1715 - val\_accuracy: 0.7185  
Epoch 4/30  
625/625 [=====] - 5s 8ms/step - loss: 6.1604e-04 - accuracy: 1.0000 - val\_loss: 2.2206 - val\_accuracy: 0.7185  
Epoch 5/30  
625/625 [=====] - 5s 8ms/step - loss: 5.1373e-04 - accuracy: 1.0000 - val\_loss: 2.2542 - val\_accuracy: 0.7185  
Epoch 6/30  
625/625 [=====] - 5s 8ms/step - loss: 4.4353e-04 - accuracy: 1.0000 - val\_loss: 2.2857 - val\_accuracy: 0.7177  
Epoch 7/30  
625/625 [=====] - 5s 8ms/step - loss: 3.9602e-04 - accuracy: 1.0000 - val\_loss: 2.3142 - val\_accuracy: 0.7192  
Epoch 8/30  
625/625 [=====] - 5s 8ms/step - loss: 3.5685e-04 - accuracy: 1.0000 - val\_loss: 2.3356 - val\_accuracy: 0.7188  
Epoch 9/30  
625/625 [=====] - 5s 8ms/step - loss: 3.2725e-04 - accuracy: 1.0000 - val\_loss: 2.3555 - val\_accuracy: 0.7184  
Epoch 10/30  
625/625 [=====] - 5s 8ms/step - loss: 3.0244e-04 - accuracy: 1.0000 - val\_loss: 2.3734 - val\_accuracy: 0.7186  
Epoch 11/30  
625/625 [=====] - 5s 8ms/step - loss: 2.8257e-04 - accuracy: 1.0000 - val\_loss: 2.3912 - val\_accuracy: 0.7188  
Epoch 12/30  
625/625 [=====] - 5s 8ms/step - loss: 2.6384e-04 - accuracy: 1.0000 - val\_loss: 2.4064 - val\_accuracy: 0.7182  
Epoch 13/30  
625/625 [=====] - 5s 8ms/step - loss: 2.4930e-04 - accuracy: 1.0000 - val\_loss: 2.4231 - val\_accuracy: 0.7183  
Epoch 14/30  
625/625 [=====] - 5s 8ms/step - loss: 2.3504e-04 - accuracy: 1.0000 - val\_loss: 2.4354 - val\_accuracy: 0.7180  
Epoch 15/30  
625/625 [=====] - 5s 8ms/step - loss: 2.2352e-04 - accuracy: 1.0000 - val\_loss: 2.4484 - val\_accuracy: 0.7175  
Epoch 16/30  
625/625 [=====] - 6s 9ms/step - loss: 2.1282e-04 - accuracy: 1.0000 - val\_loss: 2.4603 - val\_accuracy: 0.7179  
Epoch 17/30  
625/625 [=====] - 7s 11ms/step - loss: 2.0234e-04 - accuracy: 1.0000 - val\_loss: 2.4725 - val\_accuracy: 0.7177  
Epoch 18/30  
625/625 [=====] - 5s 8ms/step - loss: 1.9398e-04 - accuracy: 1.0000 - val\_loss: 2.4859 - val\_accuracy: 0.7176  
Epoch 19/30  
625/625 [=====] - 5s 8ms/step - loss: 1.8657e-04 - accuracy: 1.0000 - val\_loss: 2.4957 - val\_accuracy: 0.7178  
Epoch 20/30  
625/625 [=====] - 5s 8ms/step - loss: 1.7888e-04 - accuracy: 1.0000 - val\_loss: 2.5039 - val\_accuracy: 0.7177

Epoch 21/30  
625/625 [=====] - 5s 8ms/step - loss: 1.7195e-04 - accuracy: 1.0000 - val\_loss: 2.5150 - val\_accuracy: 0.7179  
Epoch 22/30  
625/625 [=====] - 5s 8ms/step - loss: 1.6615e-04 - accuracy: 1.0000 - val\_loss: 2.5228 - val\_accuracy: 0.7180  
Epoch 23/30  
625/625 [=====] - 5s 8ms/step - loss: 1.6032e-04 - accuracy: 1.0000 - val\_loss: 2.5317 - val\_accuracy: 0.7185  
Epoch 24/30  
625/625 [=====] - 5s 8ms/step - loss: 1.5480e-04 - accuracy: 1.0000 - val\_loss: 2.5409 - val\_accuracy: 0.7175  
Epoch 25/30  
625/625 [=====] - 5s 8ms/step - loss: 1.4990e-04 - accuracy: 1.0000 - val\_loss: 2.5485 - val\_accuracy: 0.7181  
Epoch 26/30  
625/625 [=====] - 5s 8ms/step - loss: 1.4504e-04 - accuracy: 1.0000 - val\_loss: 2.5576 - val\_accuracy: 0.7174  
Epoch 27/30  
625/625 [=====] - 5s 8ms/step - loss: 1.4059e-04 - accuracy: 1.0000 - val\_loss: 2.5652 - val\_accuracy: 0.7178  
Epoch 28/30  
625/625 [=====] - 5s 8ms/step - loss: 1.3663e-04 - accuracy: 1.0000 - val\_loss: 2.5728 - val\_accuracy: 0.7176  
Epoch 29/30  
625/625 [=====] - 6s 9ms/step - loss: 1.3282e-04 - accuracy: 1.0000 - val\_loss: 2.5788 - val\_accuracy: 0.7171  
Epoch 30/30  
625/625 [=====] - 5s 8ms/step - loss: 1.2922e-04 - accuracy: 1.0000 - val\_loss: 2.5868 - val\_accuracy: 0.7175

Plot Training and Validation Loss

```
In [16]: plt.figure(figsize = (10, 6))

loss_values = history.history['loss']
val_loss_values = history.history['val_loss']

epochs = range(1, len(loss_values) + 1)


plt.plot(epochs, loss_values, 'r', label = 'Training loss')
plt.plot(epochs, val_loss_values, 'b', label = 'Validation loss')

plt.title('Training and Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')

plt.legend()

fig = plt.gcf()
fig.savefig('results/CIFGAR10/no/train_val_loss.png')

plt.show()
```



Plot Training and Validation accuracy

```
In [17]: plt.clf()

plt.figure(figsize = (10, 6))

acc_values = history.history['accuracy']
val_acc_values = history.history['val_accuracy']

epochs = range(1, len(loss_values) + 1)

plt.plot(epochs, acc_values, 'r', label = 'Training accuracy')
plt.plot(epochs, val_acc_values, 'b', label = 'Validation accuracy')


plt.title('Training and Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')

plt.legend()

fig = plt.gcf()
fig.savefig('results/CIFGAR10/no/train_val_accuracy.png')

plt.show()
```

<Figure size 432x288 with 0 Axes>



Evaluate the Model

```
In [18]: test_loss, test_acc = model.evaluate(test_images, test_labels)

313/313 [=====] - 1s 4ms/step - loss: 2.7986 - accuracy: 0.7098
```

```
In [19]: print(f'Test accuracy: {test_acc * 100:.1f}%')
print(f'Test loss: {test_loss:.3f}')
```

Test accuracy: 71.0%  
Test loss: 2.799

Predicting the test data

```
In [20]: label_pred_test = model.predict(test_images)
label_pred_test_classes = np.argmax(label_pred_test, axis = 1)
label_pred_test_max_probability = np.max(label_pred_test, axis = 1)

313/313 [=====] - 1s 3ms/step
```

```
In [21]: # Reverse test_labels from categorical
test_labels = np.argmax(test_labels, axis = 1)
```

Visualize predictions

```
In [22]: cols = 8
rows = 2

fig = plt.figure(figsize = (2 * cols - 1, 3 * rows - 1))

for i in range(cols):
    for j in range(rows):
        random_index = np.random.randint(0, len(test_labels))
```

```

ax = fig.add_subplot(rows, cols, i * rows + j + 1)

ax.grid('off')
ax.axis('off')

ax.imshow(test_images[random_index, :])

pred_label = cifar10_classes[label_pred_test_classes[random_index]]
pred_probability = label_pred_test_max_probability[random_index]

true_label = cifar10_classes[test_labels[random_index]]

ax.set_title(f'pred:{pred_label}\nscore: {pred_probability:.3}\ntrue: {true_label}')

```



Save Model and Results

```
In [23]: model.save('results/CIFGAR10/no/mnist.h5', history)
```

```
In [24]: pickle.dump({'test_accuracy': test_acc,
                    'test_loss': test_loss,
                    'history_dict': history.history},
            open("results/CIFGAR10/no/training_metrics", "wb"))
```